

# Assignment #6

HARNEET SINGH - 400110275

COMPENG 4TN4: Image Processing

March 26, 2022

# Theory

## 1. Hough Transform

For the given image, origin is assumed at the highlighted cell as shown below:

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 255 & 255 & 255 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 255 & 0 & 255 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 255 & 255 & 255 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 255 & 0 & 0 & 255 & 0 \\ 0 & 0 & 0 & 0 & 0 & 255 & 0 & 0 & 255 & 0 \\ 0 & 0 & 0 & 0 & 0 & 255 & 0 & 0 & 255 & 0 \\ 0 & 0 & 0 & 0 & 0 & 255 & 255 & 255 & 255 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

The total length of  $\rho$  is calculated by using the Pythagoras theorem i.e.  $\text{diagonal} = \sqrt{\text{rows}^2 + \text{columns}^2}$ . Based on the Hough transform algorithm, 180 (arbitrary number)  $\theta$  and  $\rho$  values are stored in a vector.

Initially, accumulator is set to 0 and for each non-zero pixel values i.e. 18 values of 255, the  $\rho$  value is calculated as:

$$\rho = x.\cos\theta - y.\sin\theta$$

for each  $\theta$  in the accumulator. Then, calculated  $\rho$  value is rounded down to the available  $\rho$  value in the accumulator matrix and the, accumulator cell is incremented for that specific  $\rho$  and  $\theta$  value.

In the end, the accuracy of the Hough Transform is dependent on the resolution in accumulator matrix. In my accumulator matrix, the max accumulator value is 4.

As shown below, 11 lines must be detected and based on my Hough transform algorithm, I receive 11  $\rho$  and  $\theta$  pairs.



Following Python code uses the same technique to compute Hough transform on all the non-zero pixels in the given image:

NOTE: Zero-padding is not applied to the given image.

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # read given images
6 img = np.array([[0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
7                 [0, 255, 255, 255, 0, 0, 0, 0, 0, 0],
8                 [0, 255, 0, 255, 0, 0, 0, 0, 0, 0],
9                 [0, 255, 255, 255, 0, 0, 0, 0, 0, 0],
10                [0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
11                [0, 0, 0, 0, 0, 255, 0, 0, 255, 0],
12                [0, 0, 0, 0, 0, 255, 0, 0, 255, 0],
13                [0, 0, 0, 0, 0, 255, 0, 0, 255, 0],
14                [0, 0, 0, 0, 0, 255, 255, 255, 255, 0],
15                [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]])
16 img_height, img_width = img.shape[:2]
17
18 # diagonal distance between the image (max rho value)
19 diagonal = np.sqrt(np.square(img_height) + np.square(img_width))
20
21 # design accumulator based on number of theta (180) and number of rhos (180) we need
22 tick_rho = 2 * diagonal/180
23 tick_theta = 1
24
25 rhos = np.arange(-diagonal, diagonal, tick_rho)
26 thetas = np.arange(0, 180, tick_theta)
27 # print(rhos)
28 # print(thetas)
29
30 cosTheta = np.cos(np.deg2rad(thetas))
31 sinTheta = np.sin(np.deg2rad(thetas))
32
33 # defining an accumulator of size 180x180
34 accumulator = np.zeros((len(rhos), len(thetas)))
35 # print(accumulator.shape)
36
37 for i in range(img_height):
38     for j in range(img_width):
39         if img[i, j] != 0:
40             # image origin is in the middle
41             # print(f'image index is {i - img_height / 2}, {j - img_width / 2}')
42
43             # need lists to store rho and theta values
44             rhoList, thetaList = [], []
45
46             for theta in thetas:
47                 rho = ((j - img_width / 2) * (cosTheta[theta])) + ((i - img_height / 2) * (sinTheta[theta]))
48                 # print(rho, theta)
49                 rhoList.append(rho)
50                 thetaList.append(theta)
51
52             # now we need to find the index of row where we need to increment the accumulator
53
54             rhoIndex = np.argmin(np.abs(rhos - rho))
55             # print(f'{rhoIndex}')
56
57             accumulator[rhoIndex, theta] += 1
```

```

58
59 accumulatorMax = np.amax(accumulator)
60 print(f'max value in accumulator is: {accumulatorMax}')
61
62 # now we need to iterate through the accumulator matrix to find the max value
63 numMaxValues = 0
64 for k in range(len(rhos)):
65     for l in range(len(thetas)):
66         if accumulator[k, l] == accumulatorMax:
67             rho = rhos[k]
68             theta = thetas[l]
69             print(f'k, {l}, rhos is {rho}, and theta is {theta}')
70             numMaxValues += 1
71
72 print(f'number of max values in the accumulator are {numMaxValues}')

```

Result:

Line Number	$\rho$	$\theta$
1	-3	179
2	0	0
3	0	1
4	0	135
5	0	179
6	1	135
7	3	0
8	3	1
9	3	89
10	3	90
11	3	91

## 2. Optical Flow

**Zero-padding is applied to the given image to get accurate results at the edges.**

Applying kernel of length 2 to get the gradient values on the given image. For  $I_x$  considering right-pixel to be the anchor pixel.

$$I_x = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 63 & 63 & 0 & -63 & -63 \\ 0 & 127 & 63 & 0 & -63 & -127 \\ 0 & 63 & 0 & 0 & 0 & -63 \end{bmatrix}$$

For  $I_y$  considering bottom-pixel to be the anchor pixel.

$$I_y = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 63 & 191 & 255 & 191 & 63 \\ 0 & 0 & -63 & -127 & -63 & 0 \\ 0 & -63 & -127 & -127 & -127 & -63 \end{bmatrix}$$

For  $I_t$  considering 'next frame' to be the anchor pixel.

$$I_t = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & -63 & -63 & 0 & 63 & 63 \\ 0 & -127 & -63 & 0 & 63 & 127 \\ 0 & -63 & 0 & 0 & 0 & 63 \end{bmatrix}$$

Applying the flow equation and constraint equation, we can find  $\vec{u} = [u \ v]'$  because we get a over-constrained linear equation, 1 equation for each pixel value:

$A \cdot d = b$ , where A is  $9 \times 1$ , d is  $2 \times 1$  and b is  $9 \times 1$  because the spatial coherency is  $3 \times 3$  window i.e. 9 pixels in the window.

Using the given equation:  $A'A = \sum \begin{bmatrix} I_x^2 & I_x \cdot I_y \\ I_x \cdot I_y & I_y^2 \end{bmatrix}$  for each pixel using the values.

$$A'T = - \sum \begin{bmatrix} I_x \cdot I_t \\ I_y \cdot I_t \end{bmatrix}$$

Result: Vertical is rows and horizontal are columns. First value in the cell shows the horizontal velocity and second value is the vertical value. Zero-padding is applied on the given frames:

-	1	2	3	4	5	6
1	0.5, 0.5	1, 0	1, 0	1, 0	1, 0	1, 0
2	1, 0	1, 0	1, 0	1, 0	1, 0	1, 0
3	1, 0	1, 0	1, 0	1, 0	1, 0	1, 0
4	1, 0	1, 0	1, 0	1, 0	1, 0	1, 0

NOTE: values have been rounded down, please look at the supplied code for exact values.

As we can see, pixel at location (1, 1) has components in both direction but all the other ones are showing flow in the horizontal direction. This result is consistent with the zero-padding around the edges.

Following Python code uses the same technique to compute optical flow for all the pixels in the given image using  $3 \times 3$  window:

NOTE: Zero-padding is applied to the given image.

```

1  import cv2
2  import numpy as np
3  import matplotlib.pyplot as plt
4
5  # read given images
6  img0 = np.array([[0, 0, 0, 0, 0, 0],
7                  [0, 255, 255, 255, 0, 0],
8                  [0, 255, 0, 255, 0, 0],
9                  [0, 0, 0, 0, 0, 0]])
10
11 img1 = np.array([[0, 0, 0, 0, 0, 0],
12                 [0, 0, 255, 255, 255, 0],
13                 [0, 0, 255, 0, 255, 0],
14                 [0, 0, 0, 0, 0, 0]])
15 imgHeight, imgWidth = img0.shape[:2]
```

```

16
17 # zero padding
18 frame0 = np.zeros((imgHeight+2, imgWidth+2), dtype='uint8')
19 frame1 = np.zeros((imgHeight+2, imgWidth+2), dtype='uint8')
20 frame0[1:5, 1:7] = img0
21 frame1[1:5, 1:7] = img1
22
23 Grad_Ix = np.zeros_like(frame0, dtype='int')
24 Grad_Iy = np.zeros_like(frame0, dtype='int')
25 Grad_It = np.zeros_like(frame0, dtype='int')
26
27 # find partial derivative in x-direction
28 for h in range(1, imgHeight+1):
29     for k in range(1, imgWidth+1):
30         # I_x = (1/4) * (frame1[h, k] + frame1[h+1, k] + frame0[h, k] + frame0[h+1, k])
31         # I_xPlus1 = (1/4) * (frame1[frame1[h, k+1] + frame1[h+1, k+1] + frame0[h, k+1] + frame0[h+1, k+1]])
32         x_frame0 = np.sum(frame0[h:h+2, k:k+1])
33         x_frame1 = np.sum(frame1[h:h+2, k:k+1])
34
35         xPlus1_frame0 = np.sum(frame0[h:h+2, k+1:k+2])
36         xPlus1_frame1 = np.sum(frame1[h:h+2, k+1:k+2])
37
38         xPlus1 = (xPlus1_frame0 + xPlus1_frame1) * (1/4)
39         xPlus0 = (x_frame0 + x_frame1) * (1/4)
40
41         Grad_Ix[h+1, k+1] = xPlus1 - xPlus0
42
43
44 # find partial derivative in y-direction
45 for h in range(1, imgHeight+1):
46     for k in range(1, imgWidth+1):
47         y_frame0 = np.sum(frame0[h:h+1, k:k+2])
48         y_frame1 = np.sum(frame1[h:h+1, k:k+2])
49
50         yPlus1_frame0 = np.sum(frame0[h+1:h+2, k:k+2])
51         yPlus1_frame1 = np.sum(frame1[h+1:h+2, k:k+2])
52
53         yPlus1 = (yPlus1_frame0 + yPlus1_frame1) * (1/4)
54         yPlus0 = (y_frame0 + y_frame1) * (1/4)
55
56         Grad_Iy[h+1, k+1] = yPlus1 - yPlus0
57
58 # find partial derivative in t-direction
59 for h in range(1, imgHeight+1):
60     for k in range(1, imgWidth+1):
61         t_frame0 = np.sum(frame0[h:h+2, k:k+2], dtype='int')
62         t_frame1 = np.sum(frame1[h:h+2, k:k+2], dtype='int')
63
64         t = (t_frame1 - t_frame0) * (1/4)
65
66         Grad_It[h+1, k+1] = t
67
68 Ix = Grad_Ix[1:5, 1:7]
69 Iy = Grad_Iy[1:5, 1:7]
70 It = Grad_It[1:5, 1:7]
71
72 print(f'Ix is: \n {Ix}')
73 print(f'Iy is: \n {Iy}')
74 print(f'It is: \n {It}')
75
76 # Ix_square = np.square(Ix)
77 # Iy_square = np.square(Iy)

```

```

78 # I_xy = np.multiply(Ix, Iy)
79 # I_xt = np.multiply(Ix, It)
80 # I_yt = np.multiply(Iy, It)
81
82 # A_transpose_A = np.array((2, 2), dtype='int')
83 # A_transpose_b = np.array((2, 1), dtype='int')
84 # # applying least square method to find u-vector for each pixel
85 # for h in range(0, imgHeight):
86 #     for k in range(0, imgWidth):
87 #         A_transpose_A = [[Ix_square[h, k], I_xy[h, k]],
88 #                         [I_xy[h, k], Iy_square[h, k]]]
89 #
90 #         A_transpose_b = [[-I_xt],
91 #                         [-I_yt]]
92 #
93 #         u = np.linalg.lstsq(A_transpose_A, A_transpose_b, rcond=None)
94
95 A = np.zeros((9, 2), dtype='int')
96 b = np.zeros((9, 1), dtype='int')
97
98 index = 0
99 # applying least square method to find u-vector for each pixel
100 for h in range(0, imgHeight):
101     for k in range(0, imgWidth):
102         i_iterator = 0
103
104         subMatrixIx = Grad_Ix[h:h+3, k:k+3]
105         subMatrixIy = Grad_Iy[h:h+3, k:k+3]
106         subMatrixIt = Grad_It[h:h+3, k:k+3]
107
108         for i in range(3):
109             for j in range(3):
110                 A[i_iterator] = [subMatrixIx[i, j], subMatrixIy[i, j]]
111                 b[i_iterator] = -subMatrixIt[i, j]
112                 i_iterator += 1
113
114         u = np.linalg.lstsq(A, b, rcond=None)[0]
115
116         print(f'At index: {index} ({h}, {k}): u is:\n{u}', end='\n\n')
117         index += 1

```

### 3. Orientation Detection

We can perform Hough transform to get the distance of the titled box and the angles at which they are titled.



- First, load the image to perform pre-processing steps.
- Next, apply a pre-processing step to convert the image to the binary scale i.e. invert the image.
- After binarizing the image, apply Hough transform to get  $(\rho, \theta)$  values for different lines in the image.
- Next, using openCV's *hough\_line\_peaks* method(), find the peak  $(\rho, \theta)$  corresponding to lines in the image.
- With the peak values, we can explore the theta values for each line and compare the theta values. By randomly selecting, unique two  $\theta$  values, we can retrieve the difference in their alignment i.e. it will give us the orientation of the object, because  $\theta$  is the normal vector on the lines. Subtract the two  $\theta$  values to get the orientation of the object based on the lines.
- If selected properly, we can get the orientation correct in the first attempt. However, if the answer turns out to be zero, we need to replace  $\theta$  with the other values from the peaks.



# Implementation

## 1. Tracking

Python Script:

```
1  import cv2
2  import numpy as np
3  import matplotlib.pyplot as plt
4
5  # read given video
6  cap = cv2.VideoCapture('video1.mp4')
7  _, frame = cap.read()
8  oldGrayFrame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
9  frameHeight, frameWidth = oldGrayFrame.shape
10
11 numberFrames = 0
12
13 # params for ShiTomasi corner detection
14 feature_params = dict(maxCorners=50,
15                       qualityLevel=0.3,
16                       minDistance=7,
17                       blockSize=7)
18
19 # Parameters for lucas kanade optical flow
20 lk_params = dict(winSize=(25, 25),
21                 maxLevel=2,
22                 criteria=(cv2.TERM_CRITERIA_EPS | cv2.TERM_CRITERIA_COUNT, 10, 0.03))
23
24 mask = np.zeros_like(oldGrayFrame)
25 cv2.rectangle(mask, (frameWidth // 2, frameHeight // 2), (frameWidth, frameHeight), 255, -1)
26 # cv2.imshow("mask", mask)
27 # cv2.waitKey(0)
28
29 oldPts = cv2.goodFeaturesToTrack(oldGrayFrame, mask=mask, **feature_params)
30 mask = np.zeros_like(frame)
31 # cv2.imshow("newMask", mask)
32 # cv2.waitKey(0)
33
34 thickness = 35
35
36 while True:
37     ret, frame = cap.read()
38     if not ret:
39         print(f'Reached end of frames')
40         break
41
42     print(f'{numberFrames=}')
43
44     frameGray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
45
46     # optical flow
```


```

47 newPts, status, error = cv2.calcOpticalFlowPyrLK(oldGrayFrame, frameGray, oldPts, None, **lk_params)
48
49 # Select good points
50 good_new = newPts[status == 1]
51 good_old = oldPts[status == 1]
52
53 thickness = 0.95
54 # draw the tracks
55 for i, (new, old) in enumerate(zip(good_new, good_old)):
56     a, b = new.ravel().astype(np.int32)
57     c, d = old.ravel().astype(np.int32)
58     print(f'{i=}, {a=}, {b=}, {c=}, {d=}')
59     mask = cv2.line(mask, (a, b), (c, d), (0, 0, 255), int(thickness*0.5))
60     # mask = cv2.circle(mask, (a, b), int((thickness*0.5)), (0, 0, 255), -1)
61     # frame = cv2.circle(frame, (a, b), 5, (255, 0, 0), -1)
62 img = cv2.add(frame, mask)
63
64 # cv2.line(frame, (oldPts[0][0].astype(np.int32), (newPts[0][0].astype(np.int32)), (0, 0, 255), thickness + 1)
65
66 cv2.imshow("frame", img)
67
68 if numberFrames == 33:
69     cv2.imwrite("Last_Frame.jpg", img)
70     print("Image Created")
71
72 key = cv2.waitKey(1)
73 if key == 27:
74     break
75
76 numberFrames += 1
77 oldGrayFrame = frameGray.copy()
78 oldPts = good_new.reshape(-1, 1, 2)
79 # print(f'{oldPts=}')
80
81 print(f'{numberFrames=}')
82 cap.release()
83 cv2.destroyAllWindows()

```

Result:

Following image shows track of the snake by each frame, thickest point represents starting point of the snake and the thinnest point represents the last point.



## 2. MNIST Digit Recognition

Python Script:

```
1  # -*- coding: utf-8 -*-
2  """4TN4_Ass6_Impl2.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1SbmVCdUtKqRXYHfFd-SNo8pnT8zNMDjp
8  """
9
10 !pip install idx2numpy
11
12 import tensorflow as tf
13 import cv2
14 import numpy as np
15 from matplotlib import pyplot as plt
16 import idx2numpy
17 import random
18 import pandas as pd
19 from sklearn import svm, metrics
20
21 """Load dataset and convert it to numpy array."""
22
23 imgsFilePath = "train-images.idx3-ubyte"
24 labelsFilePath = "train-labels.idx1-ubyte"
25
26 images = idx2numpy.convert_from_file(imgsFilePath)
27 labels = idx2numpy.convert_from_file(labelsFilePath)
28 # print(images)
29 # print(images.shape)
30 # print(labels)
31 # print(labels.shape)
32
33 # keeping only the images with labels 0, 1 and 2:
34 images = images[labels < 3]
35 labels = labels[labels < 3]
36 # print(images)
37 # print(labels)
38 print(images.shape)
39 print(labels.shape)
40
41 """Display random 10 images"""
42
43 # print(images[0])
44 # print(labels[0])
45 # range(len(labels))
46
47 def randomTenImages(images, labels):
48     for j in range(0, 10):
49         i = random.randint(0, len(labels))
50
51         plt.imshow(images[i], cmap='gray')
52         plt.title('label: ' + str(labels[i]))
53         plt.show()
54
55 randomTenImages(images, labels)
56
57 """Split MNIST data into training, validation and test"""
58
```

```

59 numLabels = labels.shape[0]
60
61 trainSetSize = int(0.7 * numLabels)
62 ValidationSetSize = int(0.2 * numLabels)
63 testSetSize = int(0.1 * numLabels)
64 # print(trainSetSize)
65 # print(ValidationSetSize)
66 # print(testSetSize)
67
68 trainImages = images[:trainSetSize, :]
69 trainLabels = labels[:trainSetSize]
70 validationImages = images[trainSetSize:trainSetSize+ValidationSetSize, :]
71 validationLabels = labels[trainSetSize:trainSetSize+ValidationSetSize]
72 testImages = images[trainSetSize+ValidationSetSize:, :]
73 testLabels = labels[trainSetSize+ValidationSetSize:]
74 print(trainImages.shape)
75 print(trainLabels.shape)
76 # print(validationImages.shape)
77 # print(validationLabels.shape)
78 # print(testImages.shape)
79 # print(testLabels.shape)
80
81 # testing if data is split correctly (tried for all 3 splits)
82 randomTenImages(testImages, testLabels)
83
84 """## 2.1 Preprocessing: Low pass filter to remove noise and then binarize images"""
85
86 def preProcessingImgs(Images):
87     for i in range(len(Images)):
88         blurImg = cv2.GaussianBlur(Images[i], (3,3), 0)
89
90         # binarize as well
91         _, Images[i] = cv2.threshold(blurImg, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)
92
93     return Images
94
95 trainImages = preProcessingImgs(trainImages)
96 # randomTenImages(trainImages, trainLabels)
97
98 """## 2.2 Feature Extraction: Extract Contours"""
99
100 def featureExtraction(images, labels):
101     featArray = np.empty((0, 8))
102
103     for i in range(len(images)):
104         contours, _ = cv2.findContours(images[i], cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
105         cnt = contours[0]
106
107         _, _, width, height = cv2.boundingRect(cnt)
108
109         # calculate moments for the contour
110         M = cv2.moments(cnt)
111
112         if M["m00"] != 0:
113             cX = int(M["m10"] / M["m00"])
114             cY = int(M["m01"] / M["m00"])
115         else:
116             cX, cY = 0, 0
117
118
119         X_1 = cv2.countNonZero(images[i])
120         X_2 = float(width / height)

```

```

121     X_3 = cv2.contourArea(cnt)
122     X_4 = cv2.arcLength(cnt, True)
123     X_5 = cX
124     X_6 = cY
125     X_7 = width
126
127     featArray = np.append(featArray, np.array([[X_1, X_2, X_3, X_4, X_5, X_6, X_7, int(labels[i])]]), axis=0)
128
129     return featArray
130
131 featArray = featureExtraction(trainImages, trainLabels)
132 # my list has 8 elements because I am considering both moments i.e. along x and y direction,
133 # and also label is added for further processing
134 print(featArray.shape)
135 print(featArray[500])
136
137 """Display mean and variance"""
138
139 print(len(featArray))
140 # print(featArray[400][-1])
141
142 featArray_digit0 = np.empty((0, 8))
143 featArray_digit1 = np.empty((0, 8))
144 featArray_digit2 = np.empty((0, 8))
145
146 # splitting the each digits feature vectors to plot mean and variance
147 for i in range(len(featArray)):
148
149     if featArray[i][-1] == 0:
150         featArray_digit0 = np.append(featArray_digit0, [featArray[i]], axis=0)
151     elif featArray[i][-1] == 1:
152         featArray_digit1 = np.append(featArray_digit1, [featArray[i]], axis=0)
153     elif featArray[i][-1] == 2:
154         featArray_digit2 = np.append(featArray_digit2, [featArray[i]], axis=0)
155     else:
156         print("Something went wrong!!")
157
158 # print(len(featArray_digit0))
159 # print(len(featArray_digit1))
160 # print(len(featArray_digit2))
161
162 def findMeanVariance(featVector):
163
164     mean = np.mean(featVector, axis=0)
165     variance = np.var(featVector, axis=0)
166
167     # print(f'digit: {featVector[0][-1]} \n mean: {mean}, \n variance: {variance}', end='\n\n')
168
169     plt.subplot(121), plt.bar([1,2,3,4,5,6,7], mean[:7]),
170     plt.title(f"Mean of digit_{featVector[0][-1]:.0f}")
171     plt.subplot(122), plt.bar([1,2,3,4,5,6,7], variance[:7]),
172     plt.title(f"Variance of digit_{featVector[0][-1]:.0f}")
173     plt.show()
174
175     print("\n")
176
177     return mean, variance
178
179 mean_digit0, var_digit0 = findMeanVariance(featArray_digit0)
180 mean_digit1, var_digit1 = findMeanVariance(featArray_digit1)
181 mean_digit2, var_digit2 = findMeanVariance(featArray_digit2)
182

```

```

183 mean_array = np.array([mean_digit0[:7], mean_digit1[:7], mean_digit2[:7]])
184 var_array = np.array([var_digit0[:7], var_digit1[:7], var_digit2[:7]])
185
186 # print(mean_array)
187
188 meanFeatDF = pd.DataFrame(mean_array,
189                             columns = ['Non-zero Pxls', 'W / H', 'Area', 'Perimeter',
190                                         'Moment_X', 'Moment_Y', 'W'],
191                             index = ['Digit_0', 'Digit_1', 'Digit_2'])
192
193 varFeatDF = pd.DataFrame(var_array,
194                             columns = ['Non-zero Pxls', 'W / H', 'Area', 'Perimeter',
195                                         'Moment_X', 'Moment_Y', 'W'],
196                             index = ['Digit_0', 'Digit_1', 'Digit_2'])
197
198 print("\nMean Features Table:")
199 print(meanFeatDF)
200 print("\nVariance Features Table:")
201 print(varFeatDF)
202
203 """## 2.3 Dimension Reduction - Apply PCA"""
204
205 def dimensionReductionWithPlot(features, labels, numComponents):
206     mean = np.empty((0))
207     mean, eigenvectors, eigenvalues = cv2.PCACompute2(features[0:7], mean, maxComponents=numComponents)
208     # print(f"{mean}, \n{eigenvectors}, \n{eigenvalues}")
209
210     featArray7 = features[:]
211     # print(featArray7.shape)
212
213     numComponentTrainingSet = cv2.PCAProject(featArray7, mean, eigenvectors)
214
215     if numComponents == 2:
216         print(f"\n{numComponents} Components PCA:\n {numComponentTrainingSet}")
217         print(numComponentTrainingSet.shape)
218
219         plt.scatter((numComponentTrainingSet[:,0])[np.isin(labels, 0)], (numComponentTrainingSet[:,1])[np.isin(labels, 0)],
220                     color='red', marker='.')
221         plt.scatter((numComponentTrainingSet[:,0])[np.isin(labels, 1)], (numComponentTrainingSet[:,1])[np.isin(labels, 1)],
222                     color='black', marker='^')
223         plt.scatter((numComponentTrainingSet[:,0])[np.isin(labels, 2)], (numComponentTrainingSet[:,1])[np.isin(labels, 2)],
224                     color='yellow', marker='s')
225         plt.show()
226
227     return numComponentTrainingSet
228
229 twoPCAComponentTraining = dimensionReductionWithPlot(featArray, trainLabels, 2)
230
231 """## 2.4 Classification: Train SVM and plot confusion matrix"""
232
233 def svmClassifier(trainingData, labels):
234     #Create a svm Classifier
235     clf = svm.SVC(kernel='rbf', C=0.55)
236
237     #Train the model using the training sets
238     clf.fit(trainingData, labels)
239
240     #Predict the response for test dataset
241     PredictedTwoComponent = clf.predict(trainingData)
242     accuracy = metrics.accuracy_score(labels, PredictedTwoComponent)
243     # print("Accuracy:", accuracy)
244     # print(metrics.confusion_matrix(labels, PredictedTwoComponent))

```

```

245
246     return clf
247
248 twoComponentClf = svmClassifier(twoPCAComponentTraining, trainLabels)
249
250 # Predict the response for test dataset
251 # PredictedTwoComponent = clf.predict(twoPCAComponentTraining)
252 # accuracy = metrics.accuracy_score(trainLabels, PredictedTwoComponent)
253 # print("Accuracy:", accuracy)
254
255 """Confusion Matrix: with 2 PCA components"""
256
257 # metrics.confusion_matrix(trainLabels, PredictedTwoComponent)
258
259 """# 2.5 Optimizing number of components"""
260
261 # validation accuracy
262 # first need to process the validation data
263
264 validationImages = preProcessingImgs(validationImages)
265 validationFeatures = featureExtraction(validationImages, validationLabels)
266
267 # test the models for all 7 components:
268
269 for i in range(0, 7):
270     i_dimensionsFeatures = dimensionReductionWithPlot(featsArray, trainLabels, i)
271     i_ValidationFeatures = dimensionReductionWithPlot(validationFeatures, validationLabels, i)
272
273     if i == 0:
274         print(f"Training Accuracy of 7 components:\n")
275     else:
276         print(f"Training Accuracy of {i} components:\n")
277
278     clf = svmClassifier(i_dimensionsFeatures, trainLabels)
279
280     Predicted_i_Training_Component = clf.predict(i_dimensionsFeatures)
281     TrainingAccuracy = metrics.accuracy_score(trainLabels, Predicted_i_Training_Component)
282     print("Training Accuracy:", TrainingAccuracy)
283     print("Training Confusion Matrix:")
284     print(metrics.confusion_matrix(trainLabels, Predicted_i_Training_Component))
285
286
287     print(f"\nValidation Accuracy of {i} components:\n")
288     Predicted_i_Validation_Component = clf.predict(i_ValidationFeatures)
289     ValidationAccuracy = metrics.accuracy_score(validationLabels, Predicted_i_Validation_Component)
290     print("Validation Accuracy:", ValidationAccuracy)
291     print("Validation Confusion Matrix:")
292     print(metrics.confusion_matrix(validationLabels, Predicted_i_Validation_Component))
293     print("#####")
294     print(end='\n\n')
295
296 """# 2.6 Evaluation"""
297
298 # first need to process the test data
299
300 optimal_n = 2
301
302 testImages = preProcessingImgs(testImages)
303 testFeatures = featureExtraction(testImages, testLabels)
304
305 two_dimensionsFeatures = dimensionReductionWithPlot(featsArray, trainLabels, optimal_n)
306 two_testFeatures = dimensionReductionWithPlot(testFeatures, testLabels, optimal_n)

```




```

307
308 clf = svmClassifier(two_dimensionsFeatures, trainLabels)
309
310 PredictedTraining_Component = clf.predict(two_dimensionsFeatures)
311 TrainingAccuracy = metrics.accuracy_score(trainLabels, PredictedTraining_Component)
312 print("Training Accuracy:", TrainingAccuracy)
313 print("Training Confusion Matrix:")
314 print(metrics.confusion_matrix(trainLabels, PredictedTraining_Component))
315
316
317 print(f"\nTest Accuracy of {optimal_n} components:\n")
318 PredictedTest_Component = clf.predict(two_testFeatures)
319 TestAccuracy = metrics.accuracy_score(testLabels, PredictedTest_Component)
320 print("Test Accuracy:", TestAccuracy)
321 print("Test Confusion Matrix:")
322 print(metrics.confusion_matrix(testLabels, PredictedTest_Component))
323
324 """10 Random images that are correctly classified and 10 images that are incorrectly classified"""
325
326 # print(PredictedTest_Component)
327 # print(testLabels)
328
329 # print((PredictedTest_Component==testLabels).all())
330
331 comparisonList = np.equal(PredictedTest_Component, testLabels)
332 comparisonList.shape
333
334 FalseDetected = 0
335 TrueDetected = 0
336 for i, Value in enumerate(comparisonList):
337     # print(i, Value)
338     if Value == False and FalseDetected < 10:
339         # print(PredictedTest_Component[i], testLabels[i], Value)
340         print(f'Incorrect image # {FalseDetected}, false classification at index: {i}')
341         FalseDetected += 1
342
343         plt.imshow(testImages[i], cmap='gray')
344         plt.title('label: ' + str(PredictedTest_Component[i]))
345         plt.show()
346         print()
347
348     elif Value == True and TrueDetected < 10:
349         # print(PredictedTest_Component[i], testLabels[i], Value)
350         print(f'Correct image # {TrueDetected}, true classification at index: {i}')
351         TrueDetected += 1
352
353         plt.imshow(testImages[i], cmap='gray')
354         plt.title('label: ' + str(PredictedTest_Component[i]))
355         plt.show()
356         print()
357
358 if FalseDetected == 10 and TrueDetected == 10:
359     break

```


Result:


## 2 - 10 Random images with label:



### 2.1 - Pre-processing Step:


5 random images are showing after binarizing and applying gaussian blur filter:






## 2.2 - Feature Extraction:

My feature extractor has 7 features because both x and y direction moments are used.  
 For instance, features of an image: [174., 0.95, 254.5, 61.69848394, 14., 15., 19., 0., ]





#### Mean Features Table:

	Non-zero Pxls	W / H	Area	Perimeter	Moment_X	Moment_Y	\
Digit_0	162.992012	0.918013	210.915759	61.230433	13.594529	13.627693	
Digit_1	73.300524	0.429090	51.181571	46.352767	13.516021	13.418848	
Digit_2	141.578692	0.944999	107.986683	81.201382	13.361743	13.737046	

W

Digit_0	17.548293
Digit_1	8.477906
Digit_2	17.624939

#### Variance Features Table:

	Non-zero Pxls	W / H	Area	Perimeter	Moment_X	Moment_Y	\
Digit_0	736.358929	0.033930	2311.126890	124.000329	0.637578	0.629240	
Digit_1	244.770628	0.036933	231.806509	36.645001	0.932466	0.886346	
Digit_2	685.591992	0.050183	726.153576	108.739073	0.433791	0.490177	


W

Digit_0	6.922565
Digit_1	14.183334
Digit_2	6.930274

Following features: number of non zero pixels, area, perimeter and moments seem to be the most important.

## 2.3 - Dimension Reduction:

PCA with 2 components:



## 2.4 - Classification:


```
[[3809  18  304]
 [   2 4677   96]
 [   20   75 4035]]
```

Confusion matrix of the training data =

## 2.5 - Optimizing number of components:

Optimal number of components = 2 because the accuracy is the highest.

Plot below shows the accuracy of the validation data as the number of components are altered:



Training Accuracy of 1 components:

```
[[3353  18  760]
 [   0 4492  283]
 [ 344  275 3511]]
```

Training Accuracy: 0.8711261123043879

Training Confusion Matrix:

```
[[3353  18  760]
 [   0 4492  283]
 [ 344  275 3511]]
```

Validation Accuracy of 1 components:

Validation Accuracy: 0.41944146079484423

Validation Confusion Matrix:

```
[[ 257  120  810]
 [ 283 1041    0]
 [   3  946  264]]
```

#####

Training Accuracy of 2 components:

```
[[3809  18  304]
 [   2 4677   96]
 [  20   75 4035]]
```

Training Accuracy: 0.9604940165694998

Training Confusion Matrix:

```
[[3809  18  304]
 [   2 4677   96]
 [  20   75 4035]]
```

Validation Accuracy of 2 components:

Validation Accuracy: 0.825187969924812

Validation Confusion Matrix:

```
[[ 842  208  137]
 [ 155 1165    4]
 [  10  137 1066]]
```

#####

Training Accuracy of 3 components:

```
[[3827  13  291]
 [   2 4682  91]
 [  23  68 4039]]
```

Training Accuracy: 0.9625652040503222

Training Confusion Matrix:

```
[[3827  13  291]
 [   2 4682  91]
 [  23  68 4039]]
```

Validation Accuracy of 3 components:

Validation Accuracy: 0.7722878625134264

Validation Confusion Matrix:

```
[[ 894  131  162]
 [ 438  879    7]
 [   13   97 1103]]
```

#####

Training Accuracy of 4 components:

```
[[3829  12  290]
 [   2 4688  85]
 [  24  64 4042]]
```

Training Accuracy: 0.9634090211721387

Training Confusion Matrix:

```
[[3829  12  290]
 [   2 4688  85]
 [  24  64 4042]]
```

Validation Accuracy of 4 components:

Validation Accuracy: 0.7824919441460795

Validation Confusion Matrix:

```
[[ 890  125  172]
 [ 399  918    7]
 [   10   97 1106]]
```

#####

Training Accuracy of 5 components:

```
[[3829  12  290]
 [   2 4688   85]
 [  24   64 4042]]
```

Training Accuracy: 0.9634090211721387

Training Confusion Matrix:

```
[[3829  12  290]
 [   2 4688   85]
 [  24   64 4042]]
```

Validation Accuracy of 5 components:

Validation Accuracy: 0.7776584317937701

Validation Confusion Matrix:

```
[[ 891  125  171]
 [ 418   89    7]
 [  10   97 1106]]
```

#####

Training Accuracy of 6 components:

```
[[3829  12  290]
 [   2 4689   84]
 [  25   63 4042]]
```

Training Accuracy: 0.9634857318195765

Training Confusion Matrix:

```
[[3829  12  290]
 [   2 4689   84]
 [  25   63 4042]]
```

Validation Accuracy of 6 components:

Validation Accuracy: 0.7792696025778733

Validation Confusion Matrix:

```
[[ 889  125  173]
 [ 416   90    7]
 [   8   93 1112]]
```

#####



```

Training Accuracy of 7 components:

[[3829  12  290]
 [   2 4689   84]
 [  24   64 4042]]
Training Accuracy: 0.9634857318195765
Training Confusion Matrix:
[[3829  12  290]
 [   2 4689   84]
 [  24   64 4042]]

Validation Accuracy of 0 components:

Validation Accuracy: 0.7835660580021482
Validation Confusion Matrix:
[[ 886  126  175]
 [ 392  925    7]
 [   8   98 1107]]
#####

```

## 2.6 - Evaluation:

```


Training Accuracy: 0.9604940165694998
Training Confusion Matrix:
[[3809  18  304]
 [   2 4677   96]
 [  20   75 4035]]


Test Accuracy of 2 components:

Test Accuracy: 0.917337627482555
Test Confusion Matrix:
[[517  49  39]
 [  1 638   4]
 [  3  58 554]]


```

10 Correct Images:





10 Incorrect Images:



# Bibliography

- [1] 1.4. support vector machines. <https://scikit-learn.org/stable/modules/svm.html>.
- [2] Introduction to motion estimation with optical flow. <https://nanonets.com/blog/optical-flow/>.
- [3] Introduction to principal component analysis (pca). [https://docs.opencv.org/3.4/d1/dee/tutorial\\_introduction\\_to\\_pca.html](https://docs.opencv.org/3.4/d1/dee/tutorial_introduction_to_pca.html).
- [4] Optical flow. [https://opencv24-python-tutorials.readthedocs.io/en/latest/py\\_tutorials/py\\_video/py\\_lucas\\_kanade/py\\_lucas\\_kanade.html](https://opencv24-python-tutorials.readthedocs.io/en/latest/py_tutorials/py_video/py_lucas_kanade/py_lucas_kanade.html).